

## **A Neural Network Versus Black–Scholes: A Comparison of Pricing and Hedging Performances**

HENRIK AMILON\*

*Department of Economics, Lund University, Sweden*

### **ABSTRACT**

The Black–Scholes formula is a well-known model for pricing and hedging derivative securities. It relies, however, on several highly questionable assumptions. This paper examines whether a neural network (MLP) can be used to find a call option pricing formula better corresponding to market prices and the properties of the underlying asset than the Black–Scholes formula. The neural network method is applied to the out-of-sample pricing and delta-hedging of daily Swedish stock index call options from 1997 to 1999. The relevance of a hedge-analysis is stressed further in this paper. As benchmarks, the Black–Scholes model with historical and implied volatility estimates are used. Comparisons reveal that the neural network models outperform the benchmarks both in pricing and hedging performances. A moving block bootstrap is used to test the statistical significance of the results. Although the neural networks are superior, the results are sometimes insignificant at the 5% level. Copyright © 2003 John Wiley & Sons, Ltd.

**KEY WORDS** neural networks; option pricing; hedging; bootstrap; statistical inference

### **INTRODUCTION**

In the early 1970s there was a major breakthrough in financial asset modelling and option pricing theory, when Black and Scholes (1973) presented their famous formula for pricing derivatives. By modelling the price of the underlying asset,  $I_t$ , as a continuous-time diffusion

$$dI_t = \mu I_t dt + \sigma I_t dW_t \quad (1)$$

with constant parameters,  $\mu$  and  $\sigma$ , they show that a continuous rebalancing of a certain combination of the option and the underlying asset entails no risk. Hence, in the absence of arbitrage, this combination must yield the risk-free interest rate, which is also assumed to be constant. In this framework, a closed-form solution of the option price is possible. Despite the large number of questionable assumptions, the most crucial of which are the log-normality of prices, implied

---

\* Correspondence to: Henrik Amilon, Department of Economics, Lund University, Box 7082, SE-220 07, Lund, Sweden.  
E-mail: henrik.amilon@nek.lu.se

by (1), and continuous trading, the Black–Scholes prices are quite close to those observed in the market.

A great deal of effort has been made to relax some of the assumptions in the original framework, of which the stochastic volatility approach is most notable. By also letting the volatility parameter (or some transform of  $\sigma$ ) in (1) be described by a stochastic diffusion, prices are obtained with slightly better resemblance to market prices, see e.g. Scott (1987). Still, the price formulas hinge upon the continuous-time approximation, the normality of the disturbances, and the parametrization of the diffusion processes.

Hutchinson *et al.* (1994) take a different non-parametric approach. Instead of specifying the stochastic properties of the underlying diffusion, they use an artificial neural network in a non-linear regression of some input variables on the observed market prices. They do not specify how inputs affect option prices, but rather let the data determine the relationships. As inputs, they choose the price of the underlying asset normalized with the strike price of the option, and the time-to-maturity. The neural network model is tested against a benchmark, that is, the Black–Scholes formula with historical volatility estimates. The results reveal that the neural network model gives a better out-of-sample fit to observed market prices, and a lower absolute hedge error than does the benchmark model.

The practical relevance of a hedge analysis is mentioned in Hutchinson *et al.* (1994) and its importance is further stressed here. There are two ways of determining the parameters in the parametric diffusion models, one of which is to calibrate the model prices to option market data, the other to estimate the parameters from time series of the underlying asset. The parameters, and hence the model option prices, obtained in these two ways may not coincide, and this is often the case.<sup>1</sup> An obvious explanation would be that the model is misspecified, a conclusion which might be too hasty, however. How do we really know if the market option prices are correct? Believing the observed option prices to be totally wrong may seem naive, but believing all option prices to always be correct is equally naive. If they were, there would hardly be any need for any ‘improved’ option pricing model whatsoever. Another possible explanation is often referred to as the ‘peso problem’, (see Evans, 1996 for a survey of the literature on this topic). This means that the time series of the estimation and/or evaluation periods are too short, and therefore not reflect the true distribution of the underlying asset properly. Thus, the market prices incorporate the possibility of rare events, e.g. crashes, that are not represented in the sample period. Such an explanation can hardly be rejected since one may argue that the market prices reflect extreme events that have never occurred! Still the problem prevails: how do we know that the market prices are correctly adjusted for the possibility of such rare events? What can be done, and is done here, to give a fair comparison is to choose a sample period that captures the volatile states of the financial markets.

It is not surprising that calibrated model prices are closer to observed market prices, since these are used in the estimation, but they say little about the connection to the actual dynamics of the underlying asset. Regardless of the pricing model, it is, after all, the value of the underlying asset at expiration that determines the option pay-off. The market prices act as a prediction of future price variability in the underlying asset. If, on average, these predictions are not realized, then the market prices cannot be correct, and measures-of-fit based on price accuracy are of little use. A hedge analysis as described in this paper will, however, reveal if the market predictions are correct, that is, if it

<sup>1</sup> The two approaches can also be combined. Some parameters are estimated from time series data, others are calibrated to option data.

is better to fit a model to option data, or to focus on the stochastic time series properties of the underlying asset.

A neural network, on the other hand, is indeed estimated on option market data. This may seem to go against what is just mentioned, but the non-linear regression approach does not assume that all prices are always correct, but rather that a 'true' option pricing formula is embedded in the noisy market prices. If we find the true model, then we can trade on our model against the erroneous market, and hopefully make a profit. The problem is rather the specification of the relevant input variables. Do we only want to include variables from the time series of the underlying asset, or do we want to use information from the market option data, very much like the calibration approach?

In this paper I extend the non-parametric neural network approach in Hutchinson *et al.* (1994). Neural networks are estimated, or trained, on daily Swedish index call option data with additional input variables, in order to better capture the relationships between the derivative and the underlying asset. I also model the spread, that is, I use the neural network to find the mapping from the inputs to the bid and the ask prices of the options instead of assuming the mid-point, or the last traded price, to be the market price. The neural network pricing formulas are compared to two versions of a parametric benchmark model, the Black–Scholes formula with historical volatility estimates and implied volatility estimates, respectively. These benchmarks are chosen for three reasons. First, they are often used as benchmarks in comparisons with more elaborate parametric models. Second, they are frequently used by practitioners, and third, they are 'not bad'. In spite of the questionable assumptions, they are not always outperformed by more general diffusion-based models, at least in price comparisons (see e.g. Chesney and Scott, 1989).

For the reasons above, all models are compared according to both their pricing and hedging performances. Hedge analyses are also made in Hutchinson *et al.* (1994) and in Chesney and Scott (1989), but the hedging scheme used here differs from theirs in some important ways. Furthermore, by using the bootstrap technique (Efron, 1979), I perform statistical inference on my results, an issue overlooked in the earlier literature. Much of the methodology in this paper is not restricted to neural network valuation. The same steps can be taken when comparing other option pricing formulas, parametric as well as non-parametric ones.

In the second section, I give a brief description of my choice of neural network model, the Multi-Layer Perceptron (MLP). The third section describes the data, and section four more closely presents the neural network regression setup. The empirical results are given in section five. Summary and conclusions are found in the final section.

## THE MULTI-LAYER PERCEPTRON

Neurons are the basic building blocks in all artificial neural networks. In the MLP, they are organized in different layers. The inputs ( $x_k$ ) are fed to an input layer, the outputs of the network ( $o_i$ ) are given in the output layer, and in between, there is an arbitrary number of hidden layers. The hidden neurons build up an internal representation of the data. The architecture is shown in Figure 1 for an MLP with one hidden layer.

The functionality of the individual neurons is simple. Each neuron sums up the signals leading to it, adds a bias term, and makes a non-linear transformation ( $g(\cdot)$ ). The transfer function is typically a smooth monotonically increasing function, a sigmoid, such as the hyperbolic tangens or the logistic function. It can also be a linear function, but for the network to be able to map non-linear




Figure 1. A single hidden-layer MLP architecture. Notations are explained in the text

functions, some hidden layers must have non-linear transfer functions. The transformed signal of the neuron is passed on to neurons in subsequent layers, and the procedure is then repeated. The connections between the neurons are represented by weights ( $\omega_{ij}$ ,  $W_{jk}$ ). When the MLP is presented to input vectors, these are fed forward through the network via the neurons. The network outputs

$$o_i = g\left(\sum_j \omega_{ij} g\left(\sum_k W_{jk} x_k\right)\right) \quad (2)$$

are compared to known targets ( $t_i$ ) and an error function, typically the sum of squared errors

$$E = \frac{1}{2} \sum_i (t_i - o_i)^2 \quad (3)$$

is computed. The error is propagated backwards through the network and the weights are adjusted to minimize the error function. The same procedure is repeated over and over until the network outputs match the targets with an acceptable accuracy. The training algorithm in adjusting the weights is called backpropagation and was originally derived by Rumelhart *et al.* (1986).

Backpropagation finds the optimal, at least the locally optimal, weights for a given MLP architecture, but how to find the optimal architecture, that is, the optimal number of hidden layers and neurons? The issue is very much problem dependent, but it is a general statement that no more than two hidden layers are needed to approximate an arbitrary function. If the function is continuous, then only one hidden layer of sigmoid neurons is needed, (see Cybenko, 1989; Hornik, 1989). This is not necessarily the most efficient number of layers. The use of more hidden layers may result in a smaller number of neurons in total, and thereby faster training.

Regarding the number of hidden neurons, one must often, to some extent, rely on experimentation. The MLP is a universal approximator (Hornik, 1989). The implication is that given enough hidden neurons, any function can be modelled arbitrarily well. This is a very strong result, but in the presence of noise and outliers, too many neurons are likely to result in overfitting, and we end up in a situation characterized by a good fit of input vectors, but with large fluctuations in between. The network has been trained to copy the input data instead of generalizing, so when the network is presented to unseen data, its performance is low. A popular approach to avoid overfitting is to split

the data into three parts (see e.g. Hertz *et al.* 1991; Peterson and Rögnvaldsson, 1991). The training set is used for estimating the weights, the validation set is used for model selection, and the test set is used for out-of-sample evaluation. It is important to clarify that the performance of the test set must not influence the choice of architecture.

Another method for model selection is used in Anders *et al.* (1998). Drawing on the work of White (1989) and Teräsvirta *et al.* (1993), a Lagrange multiplier test procedure is used to determine whether additional hidden neurons and/or certain weight connections should be included in the architecture or not. The same idea is further explored in Rech *et al.* (2001). The main advantage of this strategy is that the model selection can be performed in-sample with no need for a special validation set.

A normalization of the input data to the same order of magnitude is also often an effective way of decreasing the complexity of the network.

## THE DATA

The data used in the empirical analysis are daily closing quotes of the Swedish OMX index and daily closing bid and ask prices of European OMX index call options, from the time periods June 1997–March 1998, and June 1998–March 1999. During this time, the world economy was highly volatile, and experienced some very large and rare fluctuations. The choice of a sample period that captures extreme events is important in dealing with the peso problem mentioned above.

The OMX index is a value weighted combination of the 30 most traded securities at the Stockholm Stock Exchange. The composition of the index is updated every sixth month, on 30 December and 30 June. On any day, OMX index options have at least three unique expiration dates; the current month, the next month, and two months ahead. On the fourth Friday of each month, some contracts expire, and new ones, with a time-to-maturity of three months but different strike prices, are introduced. The new contracts have at least three, but often five, strike prices centred around the current OMX index value. If the index moves outside the current strike price range, another strike price is added for all expiration dates to bracket that index value. Thus, the strike prices reflect the path of the OMX index during the time-to-maturity.

The reason for the *ad hoc* split into two periods is that Swedish stocks pay dividends in two months only, April and May. Since the options are of European style, options traded in April or May or expiring in these months are excluded, to avoid adjustments of the OMX index for dividend pay-offs.<sup>2</sup> After these exclusions, 514 different option contracts are traded, yielding a total of 9416 pairs of bid and ask prices.

I also use an approximation of a risk-free interest rate, which is chosen as the continuously compounded return on a 90-day treasury bill (*Statsskuldväxel*).<sup>3</sup>

## THE NEURAL NETWORK MODEL

The functional I try to model is specified to be of the form

<sup>2</sup>The sample consists of many artificial prices introduced by the option broker firm (OM), and used for 'safety calculations'. These are also excluded. The data contains no obvious outliers.

<sup>3</sup>The 90-day treasury bill is preferred to the less liquid 30- and 60-day treasury bills.

$$(c_t^b, c_t^a) = \mathbf{f}(I_t, K, T - t, T^{\text{cal}} - t, r_t, \mathbf{a})$$

where  $c_t^b$  is the bid price of the call option at time  $t$ ,  $c_t^a$  is the ask price,  $I_t$  is the OMX index value,  $K$  is the strike price,  $T - t$  is the time-to-maturity in trading days (252 per year),  $T^{\text{cal}} - t$  is the time-to-maturity in calendar days (360 per year), and  $r_t$  is the risk-free interest rate. The vector  $\mathbf{a}$  contains additional inputs. I defer the specification of  $\mathbf{a}$  until the next section, since  $\mathbf{a}$  differs depending on which of the benchmark models is used. Here, I consider only issues similar to both network models.

In order to reduce the number of inputs, I make two simplifications. First, I follow Hutchinson *et al.* (1994) and assume that  $\mathbf{f}$  is homogeneous of degree one in  $I_t$  and  $K$ , that is,  $\mathbf{f}(I_t, K, \dots) / K = \mathbf{f}(I_t/K, 1, \dots)$ . This is true for the Black–Scholes formula and it is quite a natural assumption. Otherwise, a split of the underlying asset would affect the option price. In a recent paper, Garcia and Gençay (2000) exploit the homogeneity property in terms of functional shape. Instead of using a neural network to map  $I_t/K$  and  $T - t$  directly onto the option price divided by  $K$ , they break down the pricing function into two parts. One part is controlled by  $I_t/K$  and the other by a function of  $T - t$ , similar to the functional shape of the Black–Scholes formula and other models with the same homogeneity property. They model each part with an MLP network with  $I_t/K$  and  $T - t$  as the only inputs, and show that such a restrictive model, although more time consuming to estimate, has a better pricing accuracy than the unrestrictive model by Hutchinson *et al.* This idea is not explored further here.

The second simplification concerns the fact that the time-to-maturity of a treasury bill is measured in calendar days. If  $r_t$  is only assumed to enter in the pricing function as a discount factor (as in the Black–Scholes formula), the relevant input argument is  $r_t(T^{\text{cal}} - t)$ . Thus, the specification becomes

$$(c_t^b, c_t^a)/K = \mathbf{f}(I_t/K, T - t, r_t(T^{\text{cal}} - t), \mathbf{a}).$$

The neural network model used is an MLP with one hidden layer of neurons with transfer functions  $g(\cdot) = \tanh(\cdot)$ . Since I model both the bid and the ask prices of the call options, there are two output neurons which are compared to the target vectors  $(c_t^b, c_t^a)/K$ . I choose the logistic transfer function for the output neurons, that is

$$g(\cdot) = \frac{1 + \tanh(\cdot)}{2}$$

so that the output range  $\in [0, 1]$ . Usually the output neurons are linear in function approximation, but since the targets are positive, using the logistic functions seems to make more sense.

In the earlier section, I stressed the importance of evaluating the network by the use of a training set, a validation set, and a test set. In addition, there is the trade-off between long enough estimation periods to ensure a good fit, and the risk of invoking nonstationarities, that is, changes in the relationship between input and output variables. The original periods (periods 1 and 2) are therefore split into subperiods in the following way: I train MLP networks with different numbers of hidden neurons on the first four months of period 1 (June 1997–September 1997), validate them on the consecutive two months (October 1997–November 1997), choose the one with the lowest validation error, and evaluate it on the following month (December 1997). I then repeat the procedure, but now



with the first five months (June 1997–October 1997) in the training set, the two consecutive months (November 1997–December 1997) in the validation set, and test the best network the following month (January 1998). Since period 1 ends in March 1998, this can be done another two times, thereby giving four non-overlapping test sets. Applying exactly the same procedure to period 2, we end up with a total of eight test sets (December 1997–March 1998 and December 1998–March 1999) for out-of-sample evaluation.

It is important to note that the networks can be sensitive to extrapolation. If they are presented to an input vector outside the range of the training data, their performance may be low. It is therefore of importance to assure that the validation set is 'embedded' in the training set. For example, suppose that the index drops four months and then rises the following two months. The training data, as I have chosen it, would then contain a larger amount of out-of-the money options, while the validation set is overrepresented by in-the-money options. In such a case, it is wiser to choose, say, the first, the third, the fourth, and the sixth month as training data, and use the second and the fifth month for validation, a procedure not required for the data used here.

Designing the test data in a similar way would not be correct, since the test set in a practical situation is not known in advance, and the out-of-sample evaluation should always be separated from the model selection. Instead, I monitor the largest error in the training and the validation sets. If the network gives a larger out-of-sample prediction error, that prediction is considered unacceptable and is not used.

The MLP networks are trained in batch mode, using a conjugate gradient algorithm for adjusting weights.<sup>4</sup> For each period, 20 networks are trained with three different numbers of hidden neurons. In all cases, the sum of squared errors, equation (3), is minimized for the validation set. All inputs are normalized by extracting the mean and dividing by the standard deviation of the training set.

## EMPIRICAL RESULTS

In this section, I explore how well the MLP networks do compared to the two benchmark models: the Black–Scholes formula with historically estimated volatility, and with implied volatility estimates. The implied volatility procedure means that the volatility parameter for a certain time-to-maturity is calibrated each day to get a perfect fit for an at-the-money option. The volatility estimate is then used to price other options with the same time-to-maturity.

I evaluate the competing models by using both a price and, as explained below, more importantly, a hedge analysis. The very essence of all parametric continuous-time pricing formulas is the ability to replicate an option through a dynamic hedging strategy. To be of any practical relevance, the neural network models must also be able to hedge an option position. Fortunately, Hornik *et al.* (1990), and Gallant and White (1992), show that MLP networks with a sufficient number of hidden neurons can approximate the derivative of an arbitrary non-linear function arbitrarily well. Specifically, delta-hedging is achieved by taking the derivative of (2) on the first (unnormalized) input variable. Delta-hedge analyses are also performed in Hutchinson *et al.* (1994), and in Chesney and Scott (1989). However, the hedging scheme in this paper differs substantially from theirs.

Statistical tests of the results are difficult to design properly, because of the statistical dependences of the option-price paths, as noted in Scott (1987) and Hutchinson *et al.* Here, I design such tests by using the bootstrap technique, originating from Efron (1979).

<sup>4</sup> All computations were run in MATLAB version 5.3.

### An MLP versus Black–Scholes with historical volatility

For each of the eight test periods, I train 20 single-layer MLP networks with 10, 12, and 14 hidden neurons, and choose the one with the best validation performance.<sup>5</sup> The additional input is extracted from the underlying OMX index only. I use a total of nine input variables,

$$x_t = ((I_t, I_{t-1}, I_{t-2}, I_{t-3}, I_{t-4})/K, T-t, r_t(T^{\text{cal}} - t), \hat{\sigma}_{30}, \hat{\sigma}_{10}) = (x_{t1}, \dots, x_{t9}) \quad (4)$$

which are mapped against the target vectors  $\mathbf{t}_t = (c_t^b/K, c_t^a/K)$ , yielding the network estimates  $\mathbf{o}_t = (\hat{c}_t^b/K, \hat{c}_t^a/K)$ . The volatility of the underlying asset might be suspected to be important for option pricing. The purpose of including lagged index values is to allow the neural network to model a volatility structure (or any useful distribution structure) from the data. Any number of lags can be included, but in order to economize on the input variables, only the four last lags are chosen. Previous information embedded in the time series is included by the standard deviation of the 10 and 30 most recent continuously compounded daily returns of the OMX index,  $\hat{\sigma}_{10}$  and  $\hat{\sigma}_{30}$ . Thus, the networks are presented with two slowly varying volatility trends and may respond more quickly to departures from these trends through the lagged index values.<sup>6</sup>

The dependence of the output upon the input variables is analysed by inspecting derivatives after a completed regression. For the first output, this is done by calculating

$$S_k(t) = \frac{\partial \hat{c}_t^b(x_{t1}, \dots, x_{t9})/K}{\partial x_{tk}}, \quad k = 1, \dots, 9$$

for  $t = 1, \dots, n$ , where  $n$  is the number of test vectors, and then compute

$$S_k = \frac{1}{n} \sum_t |S_k(t)| \quad (5)$$

All input variables are normalized to the same order of magnitude, so that the  $S_k$ 's are comparable in size.

The measure in (5) tells us which variables are most important. Table I shows the sensitivity dependences for the eight test periods, and for the average of the test periods. The results are valid for the first output variable, the bid price, but the dependences are similar for the second output variable, the ask price. The magnitudes of  $S_k$  sometimes differ substantially between the test periods. Averaged over all test periods, the last row, there seems to be no doubt of the most explanatory variable being  $I_t/K$  followed by  $T-t$ ,  $r_t(T^{\text{cal}} - t)$  and  $\hat{\sigma}_{30}$ . The influences of the lagged indexes and the shorter moving average standard deviation are not considered to be very important. It should be noted that these results are purely qualitative, since no statistical significance of the sensitivities are computed. Nevertheless, based on a subjective choice of the magnitudes of  $S_k$ , a re-estimation of the model could be considered.

### Pricing

The results of the out-of-sample price comparisons averaged over all test periods are shown in Table II. To investigate the performance in different parts of the input space, in-the-money options (ITM)

<sup>5</sup>In some preliminary runs, less than 10 and more than 14 hidden neurons seem to perform badly.

<sup>6</sup>The parameter estimates used in the Black–Scholes formula are  $r_t$ ,  $T^{\text{cal}} - t$ ,  $\hat{\sigma}_{90}$ , and  $T - t$ .



Table I. Sensitivity dependences, equation (5), in the MLP regressions for each testperiod, and for all eight test periods

Periods	$S_1$	$S_2$	$S_3$	$S_4$	$S_5$	$S_6$	$S_7$	$S_8$	$S_9$
1	1.26	0.06	0.08	0.09	0.07	0.36	0.32	0.07	0.04
2	1.25	0.04	0.08	0.04	0.04	0.53	0.27	0.08	0.02
3	2.08	0.14	0.09	0.05	0.03	0.31	0.53	0.18	0.22
4	2.25	0.06	0.13	0.11	0.08	0.31	0.62	0.07	0.12
5	2.16	0.10	0.11	0.06	0.10	0.69	0.39	0.25	0.03
6	2.29	0.50	0.05	0.12	0.11	0.59	0.53	0.30	0.22
7	2.41	0.08	0.05	0.06	0.24	0.66	0.45	0.85	0.09
8	1.76	0.05	0.01	0.06	0.20	0.66	0.23	0.88	0.02
1–8	1.93	0.13	0.08	0.07	0.11	0.52	0.42	0.34	0.09

and out-of-the-money (OTM) options are analysed separately. We see that the mean of the option price errors,  $ME_{BS} = n^{-1} \sum_{i=1}^n c_i - \hat{c}_i^{BS}$  and  $ME_{MLP} = n^{-1} \sum_{i=1}^n c_i - \hat{c}_i^{MLP}$ , differ between the models.<sup>7</sup> Henceforth, the subscripts are dropped unless considered to be required. The networks slightly underprice the OTM options, while Black–Scholes overprice them more heavily. The ITM options are instead slightly overpriced by the networks, while the results for Black–Scholes are more moot. Comparing all options, the network estimates show very little bias, while the Black–Scholes prices are larger than the market data. The spread is not modelled in the Black–Scholes estimates, but the Black–Scholes prices are closer to the market ask prices.

The mean error (ME) is, of course, not a good measure-of-fit, as the fluctuations around a low mean can be large. Since the network models are estimated by using the sum of squared errors, it is natural to evaluate the models with the Root Mean Squared Error (RMSE).<sup>8</sup> From Table II, we see that RMSE is lower for the MLP networks for all types of options. For the MLP networks, it is somewhat lower for the OTM options, while this is not seen for the Black–Scholes model.

Are these results statistically significant? Both Scott (1987) and Hutchinson *et al.* (1994) note that statistical tests are difficult to formulate because the errors in fitting the option prices are likely to be correlated across options and over time, but they give no guidance to how to construct such tests. Garcia and Gençay (2000) use the Diebold and Mariano (1995) test, which assumes that the sequence of the difference between the predictions of two models is covariance stationary. By estimating the autocovariance function, presumable with a suitable lag window, the null hypothesis of no difference in prediction accuracy between two models can be tested. It is important to note, however, that the prediction error series here, and also in Garcia and Gençay, are not pure time series. They consist of observations belonging to the same day as well as different days. Any particular ordering of the errors within each day is not obvious, but a reordering of the errors within each day will certainly change the estimate of the autocovariance function and hence, the Diebold and Mariano test statistic. Besides, differences in RMSE's,  $\Delta RMSE = RMSE_{BS} - RMSE_{MLP}$ , are not computable with this test.

<sup>7</sup> Due to the monitoring for extrapolation, one option in period 1 and seven options in period 5 are excluded from the MLP sample. In order to facilitate the comparisons, the same options are excluded from the Black–Scholes sample. The effects of these exclusions are in this case negligible.

<sup>8</sup> The measures in price units are obtained by using the root.

Table II. Bid and ask price error comparisons between MLP networks and Black–Scholes with historical volatility

Options	Statistics	MLP (bid)	BS (bid)	MLP (ask)	BS (ask)
OTM	ME	0.13	−2.64	0.17	−1.28
	RMSE	2.14	5.47	2.30	4.72
	$\Delta$ RMSE	3.33 (1.54,4.92)		2.42 (1.43,3.37)	
ITM	ME	−0.26	−2.34	−0.35	0.29
	RMSE	2.82	5.39	3.00	4.64
	$\Delta$ RMSE	2.57 (0.96,3.71)		1.64 (1.05,2.47)	
ALL	ME	0.01	−2.55	0.01	−0.81
	RMSE	2.36	5.45	2.53	4.69
	$\Delta$ RMSE	3.08 (1.35,4.51)		2.16 (1.39,2.84)	

Note: ME is the mean error, RMSE is the root mean squared error, and  $\Delta$ RMSE = RMSE<sub>BS</sub> − RMSE<sub>MLP</sub>. The comparisons are for 2113 OTM options, 906 ITM options, and for all 3019 options. Numbers in parantheses are bootstrapped 95% confidence intervals

Instead, I use the overlapping block resampling in Künsch (1989). As shown in Fitzenberger (1998), this method is valid under weaker assumption than covariance stationarity. The resampling scheme looks as follows. Draw a day with probability  $1/D$  from all  $D = 156$  days in the test series. Construct a block of model errors of size  $l$ , by joining the model errors for that day and the  $l - 1$  consecutive days. Repeat the procedure until the bootstrapped series consist of about the same number of blocks as the original series, and compute the statistics of interest, giving one bootstrap estimate. Calculate estimates for 1000 bootstrap replicates and extract the confidence intervals from the resampling distribution.<sup>9</sup>

The choice of the block length  $l$  is an intricate issue. Clearly, large blocks will more faithfully preserve the dependence structure in the original series. On the other hand, the resampled series must vary enough to provide a good estimate of the distribution of the test quantities, which points towards small  $l$ . Hall *et al.* (1995) give some asymptotic results for guidance. Suppose we want to estimate some statistical quantity  $\kappa$  based on a series of length  $n$ . Hall *et al.* show that under suitable conditions and for large  $n$  and  $l$ , the mean squared error of the bootstrap estimate  $\hat{\kappa}(n, l)$  is minimized for  $l \propto n^{\frac{1}{c+2}}$ . The positive integer  $c$  takes on different values depending on the statistical quantity  $\kappa$ , thereby ensuring that the number of blocks as well as the block length go to infinity as  $n$  goes to infinity. The constant of proportionality is not known, but Hall *et al.* argue that it can be estimated by comparing  $\hat{\kappa}(n, l)$  to the bootstrap estimates from a subset of shorter series of length  $m < n$ , and different block sizes  $k < l$ . Once the optimal block size  $\hat{k}$  for a series of length  $m$  is determined, the block size for a series of length  $n$  is given by

$$\hat{l} = \hat{k} \left( \frac{n}{m} \right)^{\frac{1}{c+2}}$$

<sup>9</sup>When bootstrapping the differences of two models, the same random draws are applied to each series.




Figure 2. The width of the block bootstrapped confidence intervals for the difference in RMSE for all 3019 bid option price errors, plotted versus the block length  $l$

This procedure has certain drawbacks. First, it is obviously very computer intensive, since it suggests bootstrapping a large number of subseries for each  $k$ . Second, the results are valid asymptotically, that is,  $n$  and  $m$  are large, and block sizes  $k$  and  $l$  are large, particularly large enough to capture all serial dependences in the data. Third, as exemplified in Davison and Hinkley (1997), different choices of  $m$  will result in different  $\hat{k}$  and hence, different  $\hat{l}$ . At best, the above procedure will result in a range of  $\hat{l}$  from which the bootstrap estimate is extracted, relying to some extent on a subjective choice.

For the above reasons, a more heuristic approach is used in choosing  $l$ , when determining the 95% confidence interval for  $\Delta\text{RMSE}$ . If  $l = 1$ , dependences between days are not accounted for. But if there are interdaily dependences and  $l$  is increased, the confidence interval should widen until the block length is large enough to capture the dependences. Eventually, the width of the confidence interval falls because of the small variability in the resampled blocks. We are then able to identify a plateau region of  $C_{0.975,l} - C_{0.025,l}$ , where  $C_{\alpha,l}$  is the bootstrapped  $\alpha$ -percentile with a block length  $l$ , and we choose the confidence interval as the means of  $C_{0.975,l}$  and  $C_{0.025,l}$ , with  $l \in$  plateau region. If anything, we are too conservative in determining our confidence intervals.

Figure 2 shows the width of the confidence interval as a function of  $l$  for  $\Delta\text{RMSE}$ , estimated from all the 3019 pairs of model bid price errors in Table II. The plateau region is identified for  $l$  in the range 20–26, yielding the confidence interval on the lower left in Table II. The same procedure is then applied to all other RMSE comparisons in the table. It can be seen that in all comparisons with

the Black–Scholes model, the RMSEs are significantly lower, at the 5% level, for the MLP networks. Although not pursued here, the same steps can be taken in determining the confidence intervals for the other estimates in Table II.

### *Hedging*

At the heart of all diffusion-based option pricing theory lies the replication of an option with a portfolio of other assets. The ability to hedge an option can be used as a measure-of-fit in model comparisons and, in many cases, it is a more accurate way to proceed. The market option prices reflect the anticipation of the future distribution of the underlying asset. If these expectations are not fulfilled on average, then the market option prices are wrong and, consequently, measures-of-fit based on price accuracy are of little use. This may seem a contradiction since the MLP networks are trained on option market data. However, it is important to clarify that in all regressions, linear or non-linear, the observations are assumed to be contaminated with noise, and the purpose is to reveal the hidden deterministic relationship.

A hedge analysis, however, would reveal which of the competitive models works best in the following way. Suppose that, at some point in time, for each model, we buy a *theoretical* option, the price of which is given by the model, and hedge the option until some future date, possibly the maturity date, when we sell the theoretical option and close our positions. The tracking error, that is, the difference between the value of the option and the hedge position, can be used as a measure of accuracy. Either the tracking error at the termination of the hedge position or some statistics based on the tracking error during the hedge can be used, but it is the absolute value of the tracking errors that is of importance; if a theoretical option is instead sold, all signs are reversed. Since all prices and hedge ratios are calculated theoretically, the model with the lowest absolute tracking error has best captured the dynamics of the underlying asset.

Hutchinson *et al.* (1994) perform a similar hedge analysis with the absolute value of the tracking error at the date of expiry as a measure-of-fit. They hedge, however, a market option, not a theoretical option, which cannot be correct. Suppose their model under scrutiny is the ‘true’ one—a change in the model option price is exactly offset by a change in the hedge. If they buy a theoretical option and hedge it to expiration, the tracking error would be zero. If the market price is larger than the theoretical price, the tracking error at expiration equals this difference and, consequently, the model is considered to be less accurate. Using the market option prices in this way is certainly illuminating, but it does not clarify if market prices correctly reflect the variability of the underlying asset.

Another way of assessing model performance, addressing the question of correctness in market prices, and with a more practical value of any improvements in pricing accuracy, is to trade in mispriced options. If at all times and for all options and all competitive models, we *buy* those options that are underpriced according to the models, that is, if the model bid prices are higher than the market ask prices, and *sell* the options that are overpriced, that is, when the model ask prices are lower than the market bid prices. We then delta-hedge each option until the overpriced options become underpriced, or the underpriced options become overpriced, or the options expire, whatever happens first. If a model can successfully identify mispriced options, then the value of the terminal position should be positive, and a model yielding a higher final value is to be preferred to one with a lower final value. Thus, the sum of the terminal profits can be used to rank different models.<sup>10</sup>

<sup>10</sup>Due to discrete rebalancing and model misspecifications, these profits are of course not without risk (not even under continuous trading). The standard deviation of the profits may then also be of interest.

Here, I do use the market options, but in a different way from Hutchinson *et al.* (1994). If the model price is a better predictor of future variability, then we make an initial profit, which is, on average, added to the wealth. If not, then the model is truly misspecified, and we will lose money in the long run.<sup>11</sup>

To be more specific, the delta-hedging used here works as follows. Let  $V_t^I$ ,  $V_t^B$ , and  $V_t^C$  denote the amount invested in the index, the risk-free asset, and the call option at time  $t$ . If the market option is underpriced at time 0, that is if  $c_0^a < \hat{c}_0^b$ , we buy the option ( $V_0^C = c_0^a$ ), sell the index ( $V_0^I = -I_0\Delta_0$ ,  $\Delta_0 = \frac{\partial \hat{c}_0^b}{\partial I_t}|_{t=0}$ ), and put the rest in the risk-free asset ( $V_0^B = I_0\Delta_0 - c_0^a$ ). When we close the position, the option must be sold at the bid price, hence the expression for  $\Delta_0$ . To illustrate how the hedge works, we can decompose  $V_0^C$  and  $V_0^B$  into

$$\begin{aligned} V_0^C &= c_0^a = (c_0^a - \hat{c}_0^b) + \hat{c}_0^b \equiv V_0^{C'} + V_0^{C''} \\ V_0^B &= I_0\Delta_0 - c_0^a = -(c_0^a - \hat{c}_0^b) + (I_0\Delta_0 - \hat{c}_0^b) \equiv V_0^{B'} + V_0^{B''} \end{aligned}$$

Prior to the maturity time  $T$ , and as long as the option is not overpriced, that is, as long as not  $\hat{c}_t^b > \hat{c}_t^a$ , the positions are daily rebalanced according to

$$\begin{aligned} V_t^I &= -I_t\Delta_t, \quad \Delta_t = \frac{\partial \hat{c}_t^b}{\partial I_t} \\ V_t^C &= c_t^b = (c_t^b - \hat{c}_t^b) + \hat{c}_t^b \equiv V_t^{C'} + V_t^{C''} \\ V_t^B &= e^r V_{t-1}^B + I_t(\Delta_t - \Delta_{t-1}) = e^r V_{t-1}^{B''} + (e^r V_{t-1}^{B'} + I_t(\Delta_t - \Delta_{t-1})) \equiv V_t^{B'} + V_t^{B''} \end{aligned}$$

Again, the option must be sold at the bid price, hence the expression for  $V_t^C$ . The value of the replicating portfolio is then

$$V_t = V_t^{C'} + V_t^{B'} + V_t^{C''} + V_t^{B''} + V_t^I = (c_t^b - \hat{c}_t^b) + e^{\sum_{u=1}^t r_u} (\hat{c}_0^b - c_0^a) + V_t^{C''} + V_t^{B''} + V_t^I$$

The last three terms only involve model bid prices and derivatives, so if we have identified the correct model and rebalance continuously, the sum of these terms should be zero. In addition, our trading strategy ensures that  $V_t^{C'} > 0$  and  $V_t^{B'} > 0$  when we close our position at  $t = \text{term}$ . Since each model is used to buy low and sell high, a successful model is expected to have a high terminal value,  $V_{\text{term}} = V_{\text{term}}^I + V_{\text{term}}^B + V_{\text{term}}^C$ , with little variation.<sup>12</sup> From this, we can define a performance measure:

$$P = \sum_i e^{-n_{0,i} \times \text{term},i} V_{\text{term},i} \quad (6)$$

which is the sum of all discounted terminal profits.

The results are in Table III, where  $P$  is shown for all test periods. The performance measure for periods 1–8 is not just an average of all periods. At the end of each test period, the positions are

<sup>11</sup>Chesney and Scott (1989) also trade in mispriced options. They buy underpriced options and sell overpriced options, but they close their hedges the next day regardless of whether the options are still mispriced or not, which is odd. Also, their gains are computed from the midpoint of the bid–ask spread and not from actual prices.

<sup>12</sup>Since  $V_t^{C'} = c_t^b - \hat{c}_t^b$ , we can terminate the hedge before maturity of the option already if  $\hat{c}_t^b > \hat{c}_t^a$  in order to have a positive (expected) terminal value.

Table III. Hedge error comparisons between MLP and Black–Scholes with historical volatility, for all eight test periods, and for the overall period

Models	Statistics	Periods								
		1	2	3	4	5	6	7	8	1–8
MLP	$P$	234	−145	−150	10	−47	−331	362	−137	2096 (−714,5180)
	Std. dev.	5.57	4.19	3.79	2.85	19.58	7.48	4.84	5.96	11.02
	No. of hedges	216	157	105	59	108	240	175	97	1180
BS	$P$	426	−685	−942	515	−2302	−960	677	352	−5019 (−14334,4537)
	Std. dev.	1.68	2.98	3.68	2.91	6.82	4.00	5.36	2.63	7.17
	No. of hedges	330	323	285	123	296	369	437	154	2440
	$\Delta P$									7116 (−458,15210)

Note:  $P$  is the performance measure, equation (6), and  $\Delta P$  is the difference in  $P$  between Black–Scholes and MLP networks. Numbers in parantheses are bootstrapped 95% confidence intervals

closed even though, for example, an underpriced option is still underpriced. For the overall period in the last column, all options are hedged as described above. This means, for instance, that I switch MLP networks between two test periods. In Table III, it is shown that in many periods, hedging results in loss of money. In particular, the fifth period causes problems for both models, with a large loss for Black–Scholes and a high standard deviation for the network. This is partly due to the fact that the hedge positions at the end of the test periods must be closed, but also because I rebalance the positions on a daily basis only. A more frequent hedging would result in a smaller dispersion of the terminal values.

To some degree, Table III illustrates the variation of the test periods, but the relevant measure of comparison is found in the last column. The accumulated (discounted) wealth from the hedges is almost 2100 for the MLP networks, compared to a loss of 5020 for the Black–Scholes model.<sup>13</sup> The worse fit to the market option data in Table II for Black–Scholes results in the identification of more mispriced options. The Black–Scholes trader takes 2440 hedge positions, while the MLP trader makes 1180 hedges. Unfortunately, the Black–Scholes trader does a poor job in finding mispriced options.

The significance of the results can be checked in a way similar to the one used in the earlier section. For each model, I construct time series by summing up the terminal profits belonging to the same terminal day. The series are extended by adding zeros for those intervening days with no termination. The moving block bootstrap is then applied to the series with various block sizes and the 95% confidence intervals for the sum of each series, as well as the difference of the sums, are calculated as described earlier.

In Table III,  $P$  is not significantly different from zero, neither for the network model nor for the Black–Scholes model. Furthermore, the difference,  $\Delta P$ , is insignificant at the 5% level.

<sup>13</sup> The profits are in Swedish currency (SEK), and calculated from a position in one option only. Since the minimum tradable amount is 100 options, the profits/losses from the trading strategies should be multiplied with 100. I have not considered the effects of trading *all* available options at the bid and ask prices.



Table IV. Bid and ask price error comparisons between MLP networks and Black–Scholes with implied volatility

Options	Statistics	MLP (bid)	BS (bid)	MLP (ask)	BS (ask)
OTM	ME	0.10	−0.59	−0.05	−0.67
	RMSE	1.00	1.36	1.22	1.48
	ΔRMSE	0.37 (−0.00,0.80)		0.26 (−0.02,0.61)	
ITM	ME	−0.48	0.03	−0.61	1.28
	RMSE	1.70	1.64	1.85	2.10
	ΔRMSE	−0.06 (−0.31,0.13)		0.25 (−0.19,0.72)	
ALL	ME	−0.07	−0.40	−0.22	−0.09
	RMSE	1.25	1.45	1.44	1.69
	ΔRMSE	0.20 (−0.08,0.54)		0.25 (−0.05,0.61)	

Note: ME is the mean error, RMSE is the root mean squared error, and  $\Delta\text{RMSE} = \text{RMSE}_{\text{BS}} - \text{RMSE}_{\text{MLP}}$ . The comparisons are for 2119 out-of-the-money options, 908 in-the-money options, and for all 3027 options. Numbers in parantheses are bootstrapped 95% confidence intervals

### An MLP versus Black–Scholes with implied volatility

The Black–Scholes formula with implied volatility is all about relative pricing. Each day, the volatility parameter is adjusted to coincide with the price of an at-the-money (ATM) option with a certain time-to-maturity.<sup>14</sup> The resulting implied volatility is used to price other options with the same date of expiry. Because of the spread of the ATM option, there will be one implied volatility for bid prices and one for ask prices. The underlying assumption is that ATM options are correctly priced, and consequently, all options should be priced relative to them.

The analogy in MLP regressions is to invoke the price of an ATM option with the same time-to-maturity in the input variables, that is

$$\mathbf{x}_t = (I_t/K, T-t, r_t(T^{\text{cal}}-t), (c_t^b/K)_{\text{ATM}}, (c_t^a/K)_{\text{ATM}}, (I_t/K)_{\text{ATM}}) = (x_{t1}, \dots, x_{t6}) \quad (7)$$

In a similar manner as above, for each of the eight periods, I train 20 single-layer MLP networks with 8, 10, and 12 hidden neurons, choose the one with the lowest validation error, and use it for out-of-sample pricing and hedging.

### Pricing

The price comparisons are shown in Table IV. It is not surprising that the measures-of-fit are better than in Table II, since I now use information from a subset of the options I am pricing.<sup>15</sup> From the means of the model errors, it appears that the networks slightly overprice the (fewer) ITM options, while Black–Scholes underprices them, at least for the ask prices. For the OTM options, the network prices have little bias, while the Black–Scholes prices are higher than the market prices. It seems as

<sup>14</sup>The ATM options are, for each maturity, chosen as the ones with  $I_t/K$  closest to one.

<sup>15</sup>The monitoring for extrapolation no longer results in any exclusions of the MLP price predictions.

Table V. Hedge error comparisons between MLP and Black–Scholes with implied volatility, for all eight test periods, and for the overall period

Models	Statistics	Periods								
		1	2	3	4	5	6	7	8	1–8
MLP	<i>P</i>	96	39	−43	−171	−405	61	−145	−73	−266 (−1918,1381)
	Std. dev.	4.82	5.41	4.29	3.83	15.76	3.54	7.27	5.03	12.13
	No. of hedges	87	98	31	42	77	48	56	80	548
BS	<i>P</i>	23	−4	−118	−26	−709	129	−282	115	−1766 (−3300,−379)
	Std. dev.	1.59	1.93	5.80	5.14	8.62	4.13	7.77	2.51	7.46
	No. of hedges	76	65	118	60	128	99	78	56	707
	$\Delta P$									1500 (−271,3376)

Note: *P* is the performance measure, equation (6), and  $\Delta P$  is the difference in *P* between Black–Scholes and MLP networks. Numbers in parantheses are bootstrapped 95% confidence intervals

if the well-known ‘volatility smile’ is somewhat skewed in this market. The comparison of all options indicates a small overpricing for both models.

The root mean squared error is lower for the network models, except for the ITM bid prices. However, for the overall comparison, RMSE is lower for the MLP networks. Using the same bootstrap method to determine the 95% confidence intervals as above, we find that the differences in RMSE’s are insignificant, although the differences for the OTM options are only slightly so, in favor of the neural network models. The plateau region mentioned above is in this particular case identified for *l* in the range of 12–17.

### Hedging

Our competitive models no longer have a clear connection to the time series properties of the underlying index. To a large extent, the results rely on the accuracy of the ATM options. If the expectations of future variability in the index, embedded in the market price, do not come true, then some other models might perform better, despite a worse fit to the market prices. In particular, the hedge performances should not only be used to compare the two models in this section, but also to compare the models in the previous section.

The results from the delta-hedge schemes are shown in Table V.<sup>16</sup> The variation of *P* is seen over the eight test periods. Once more, period five seems to be a difficult period for both models, with large losses and high sample standard deviations. Because of the effects of forcing the hedges to close at the end of each test period, a fair comparison is given in the last column only. Trading in both models results in losses, but the loss for the MLP trader is limited to 266, compared to 1767 for the Black–Scholes trader. The number of hedges are 548 for the MLP trader compared to 707 for the Black–Scholes trader. The larger number of hedges for the Black–Scholes model is reflected in the worse fit to market data in Table IV.

<sup>16</sup>Since the ATM options are included in the regression, equation (7), the hedge may be improved by a suitable position in these options.

Table VI. Differences in hedge error comparisons between MLPh and the two models, MLPi and BSi

Models	$\Delta P$
MLPh-MLPi	2363 (-1759, 6355)
MLPh-BSi	3863 (291,7764)

*Note:*  $\Delta P$  is the difference in  $P$ , equation (6), between Black–Scholes and MLP networks. Numbers in parantheses are bootstrapped 95% confidence intervals

The bootstrapped 95% confidence intervals reveal that the Black–Scholes model have a  $P$  significantly less than zero. Nevertheless, the difference,  $\Delta P$ , is insignificant at this level.

Finally, Table VI shows the differences in  $P$  for the neural network model that uses information from the time series of the index (henceforth MLPh) and the two models in this section (MLPi and BSi). The two neural network models do not perform significantly different, while MLPh is significantly better than the Black–Scholes model with implied volatility estimates, at the 5% level.

## SUMMARY AND CONCLUSIONS

In this paper, I train feedforward neural networks, more specifically MLP networks, in a non-linear regression of some input variables on daily Swedish call option data. I use both the bid and ask prices in order to model the spread, instead of the common approach of assuming the midpoint to be the market price. I focus on two regression set-ups. In the first regression, I use inputs solely from the time series properties of the underlying asset (MLPh), while in the other, I infer information from the observed option prices (MLPi). The neural network pricing formulas are compared to two benchmarks, the Black–Scholes formula with historical volatility estimates (BSh), and with volatility estimated from observed ATM options (BSi), respectively. All models are compared both according to their pricing accuracy and, as argued, more importantly, to their hedging performances. Unfortunately, the complicated dependences of the option price paths make statistical inference difficult. Here, a moving block bootstrap is used to determine the statistical significance of the results.

From a pricing perspective, I find that MLPi has the lowest RMSE, followed by BSi, MLPh, and BSh. On the other hand, using the models to trade and hedge mispriced options, I find that MLPh is preferred to MLPi, followed by BSi and BSh. The results indicate that although pricing options relative to observed ATM options, as in MLPi and BSi, gives a better fit to market data, this may not be the best way to proceed. Relying on the connection between the options and the underlying asset, as modeled in MLPh, gives a better hedge. When using a hedge analysis to compare models, I find that MLPh is significantly better than BSi at the 5% level, but I cannot, for example, conclude that MLPh is better than BSh, even if it is likely.

An important result from a market efficiency perspective is that I cannot say that market prices are wrong, since no model gives a significant profit, at least not when the hedge positions are rebalanced on a daily basis. Trading on BSi gives a significant loss, however. This could imply that ATM

options are wrong, but MLPi, also using ATM options, gives no significant loss. Hence, the conclusion would rather be that BSi does not give the proper mapping from the ATM options to other options with the same time-to-maturity.

The overall conclusion is that although the neural network models are superior, the results are often insignificant at the 5% level. Furthermore, the network models outperform the Black–Scholes formula, but the results may alter when compared to more elaborate parametric pricing models.

In practice, daily hedging is not very realistic, but I see no reason why the same methods cannot be used for data with higher frequency, or ultimately, tick data. With irregular sampling intervals, the time between trades must be included in the regressions where time series properties from the underlying asset are used.

## ACKNOWLEDGMENTS

The author is grateful to Ola Hössjer and Tobias Rydén for valuable help on bootstrapping. Also thanks to Christine Brown for insightful overall comments on the paper.

## REFERENCES

- Anders U, Korn O, Schmitt C. 1998. Improving the pricing of options: a neural network approach. *Journal of Forecasting* **17**: 369–388.
- Black F, Scholes M. 1973. The pricing of options and corporate liabilities. *Journal of Political Economy* **81**: 637–654.
- Chesney M, Scott L. 1989. Pricing european currency options: a comparison of the modified Black–Scholes model and the random variance model. *Journal of Financial and Quantitative Analysis* **24**: 267–284.
- Cybenko G. 1989. Approximation by superpositions of a sigmoidal function. *Mathematics of Control, Signals and Systems* **2**: 303–314.
- Davison A, Hinkley D. 1997. *Bootstrap Methods and their Application*. Cambridge University Press: Cambridge.
- Diebold F, Mariano R. 1995. Comparing predictive accuracy. *Journal of Business and Economic Statistics* **13**: 253–263.
- Efron B. 1979. Bootstrap methods: another look at the jackknife. *Annals of Statistics* **7**: 1–26.
- Evans M. 1996. Peso problems: their theoretical and empirical implications. In *Handbook of Statistics: Statistical Models in Finance*. Maddala G, Rao CR (eds). North-Holland: Amsterdam.
- Fitzenberger B. 1998. The moving blocks bootstrap and robust inference for linear least squares and quantile regressions. *Journal of Econometrics* **82**: 235–287.
- Fletcher R. 1990. *Practical Methods of Optimization*, 2nd edn. John Wiley: New York.
- Gallant A, White H. 1992. On learning the derivatives of an unknown mapping with multilayer feedforward networks. *Neural Networks* **5**: 128–138.
- Garcia R, Gençay R. 2000. Pricing and hedging derivative securities with neural networks and a homogeneity hint. *Journal of Econometrics* **94**: 93–115.
- Hall P, Horowitz J, Jing B. 1995. On blocking rules for the bootstrap with dependent data. *Biometrika* **82**: 561–574.
- Hertz J, Krogh A, Palmer R. 1991. *Introduction to the Theory of Neural Computation*. Addison-Wesley: Redwood City, CA.
- Hornik K. 1989. Multilayer feedforward networks are universal approximators. *Neural Networks* **2**: 359–366.
- Hornik K, Stinchcombe M, White H. 1990. Universal approximation of an unknown mapping and its derivatives. *Neural Networks* **3**: 551–560.
- Hutchinson J, Lo A, Poggio T. 1994. A nonparametric approach to pricing and hedging derivative securities via learning networks. *Journal of Finance* **49**: 851–889.
- Künsch H. 1989. The jackknife and bootstrap for general stationary observations. *Annals of Statistics* **17**: 1217–1241.

- Peterson C, Rögnvaldsson T. 1991. *An Introduction to Artificial Neural Networks*. Department of Theoretical Physics. Lund University: Lund.
- Reh G, Teräsvirta T, Tschernig R. 2001. A simple variable selection technique for nonlinear models. *Communications in Statistics, Theory and Methods* **30**: 1227–1241.
- Rumelhart D, Hinton G, Williams R. 1986. Learning internal representation by error propagation. In *Parallel Distributed Processing: Explorations in the Microstructure of Cognition, Volume 1: Foundations*, Rumelhart D, McClelland J (eds). MIT Press: Cambridge, MA.
- Scott L. 1987. Option pricing when the variance changes randomly: theory, estimation, and an application. *Journal of Financial and Quantitative Analysis* **22**: 419–438.
- Teräsvirta T, Lin C-F, Granger C. 1993. Power of the neural network linearity test. *Journal of Time Series Analysis* **14**: 209–220.
- White H. 1989. An additional unit test for neglected non-linearity in multilayer feedforward networks. *Proceedings of the International Joint Conference on Neural Networks*, San Diego: 451–455.

*Author's biography:*

**Henrik Amilon** completed an MSc in Engineering Physics at the Lund Institute of Technology. He received his PhD in Financial Economics from Lund University. After post-doc training at the School of Finance and Economics at University of Technology, Sydney, he is currently with Nordea, Copenhagen and Stockholm.

*Author's address:*

**Henrik Amilon**, Department of Economics, Lund University, PO Box 7082, SE-220 07, Lund, Sweden.